



Bloomberg AI

Bloomberg is building the world's most trusted information network for financial professionals. Our core product, the Bloomberg Terminal, is an independent and unbiased source of information for our clients – everyone from C-Suite executives, traders, analysts, government officials, business heavyweights, and news professionals around the globe.

6,000+ technologists – and counting – are dedicated to advancing and building new solutions for the Bloomberg Terminal in order to solve complex, real-world problems. No other company processes the breadth and depth of financial data into meaningful and actionable information as well, or as quickly, as we do. Our engineers and researchers work together to roll out new software daily that integrates seamlessly and immediately into our clients' workflows.

Our technology expertise is changing the world and makes a positive impact in our local communities. Bloomberg engineers are active members of the open source community, and contribute to and leverage open source software in our products. And, nearly all of our company profits go directly to Bloomberg Philanthropies to support lasting change in the environment, government innovation, education, arts, and public health.

Artificial Intelligence @ Bloomberg

We are a close-knit team of 150+ researchers and engineers, focusing on projects related to AI, ML, NLP, NLU, IR, and QA.

We collaborate with researchers from academia by publishing papers, providing funding for research and fellowships, participating in conferences and program committees, hosting interns, and peer-reviewing scientific articles.

We develop solutions for semantic and syntactic parsing problems, NER & NEL systems, sentiment analysis, market impact indicators, knowledge base construction, social media analysis, question answering systems, and high performance IE. Our customers rely on these systems to make swift financial decisions; therefore, we need to guarantee precision, accuracy and latency beyond what academia and many other industries would demand.

We are looking for research scientists and engineers who are eager to apply Machine Learning and NLP to solve challenging open-ended problems. If you want to be part of a team making an impact on the global financial industry and are not afraid to get your hands dirty in data, check out the following opportunities within our Artificial Intelligence group:

Bloomberg
Engineering

Bloomberg @ ACL

Learn more about the work we're doing by speaking to any of our engineers at the conference.

Papers

In 2019, Bloomberg's engineers and data scientists published 20+ research papers on AI, Machine Learning, and NLP at conferences such as: AAAI, NAACL-HLT, SIGIR, ACL and NeurIPS. Selections follow:

- Ravneet Arora, et al., "[A Semi-Markov Structured Support Vector Machine Model for High-Precision Named Entity Recognition](#)," ACL 2019.
- Yuntian Deng, David Rosenberg, and Gideon Mann, "[Challenges in end-to-end neural scientific document OCR](#)," ICDAR 2019.
- Minjie Xu and Gary Kazantsev, "[Understanding Goal-Oriented Active Learning via Influence Functions](#)," Machine Learning with Guarantees Workshop at NeurIPS 2019.

Read more about their work and view a selection of their research at TechAtBloomberg.com/AI.

Learn more

// Email our recruiters to learn more about full-time, co-op, post-doc, and internship roles:

New York: Irinka - itoidze@bloomberg.net

London: Fran - fcucinotta@bloomberg.net

// Learn about our data science research grants & fellowships:

TechAtBloomberg.com/data-science-research-grant-program/

// Explore our Foundations of Machine Learning video course:

TechAtBloomberg.com/FOML

// Read about the unique problems we're solving:

TechAtBloomberg.com/post-topic/data-science

// Explore our open source projects: www.github.com/Bloomberg

// Follow us on Twitter: [@TechAtBloomberg](https://twitter.com/TechAtBloomberg)

